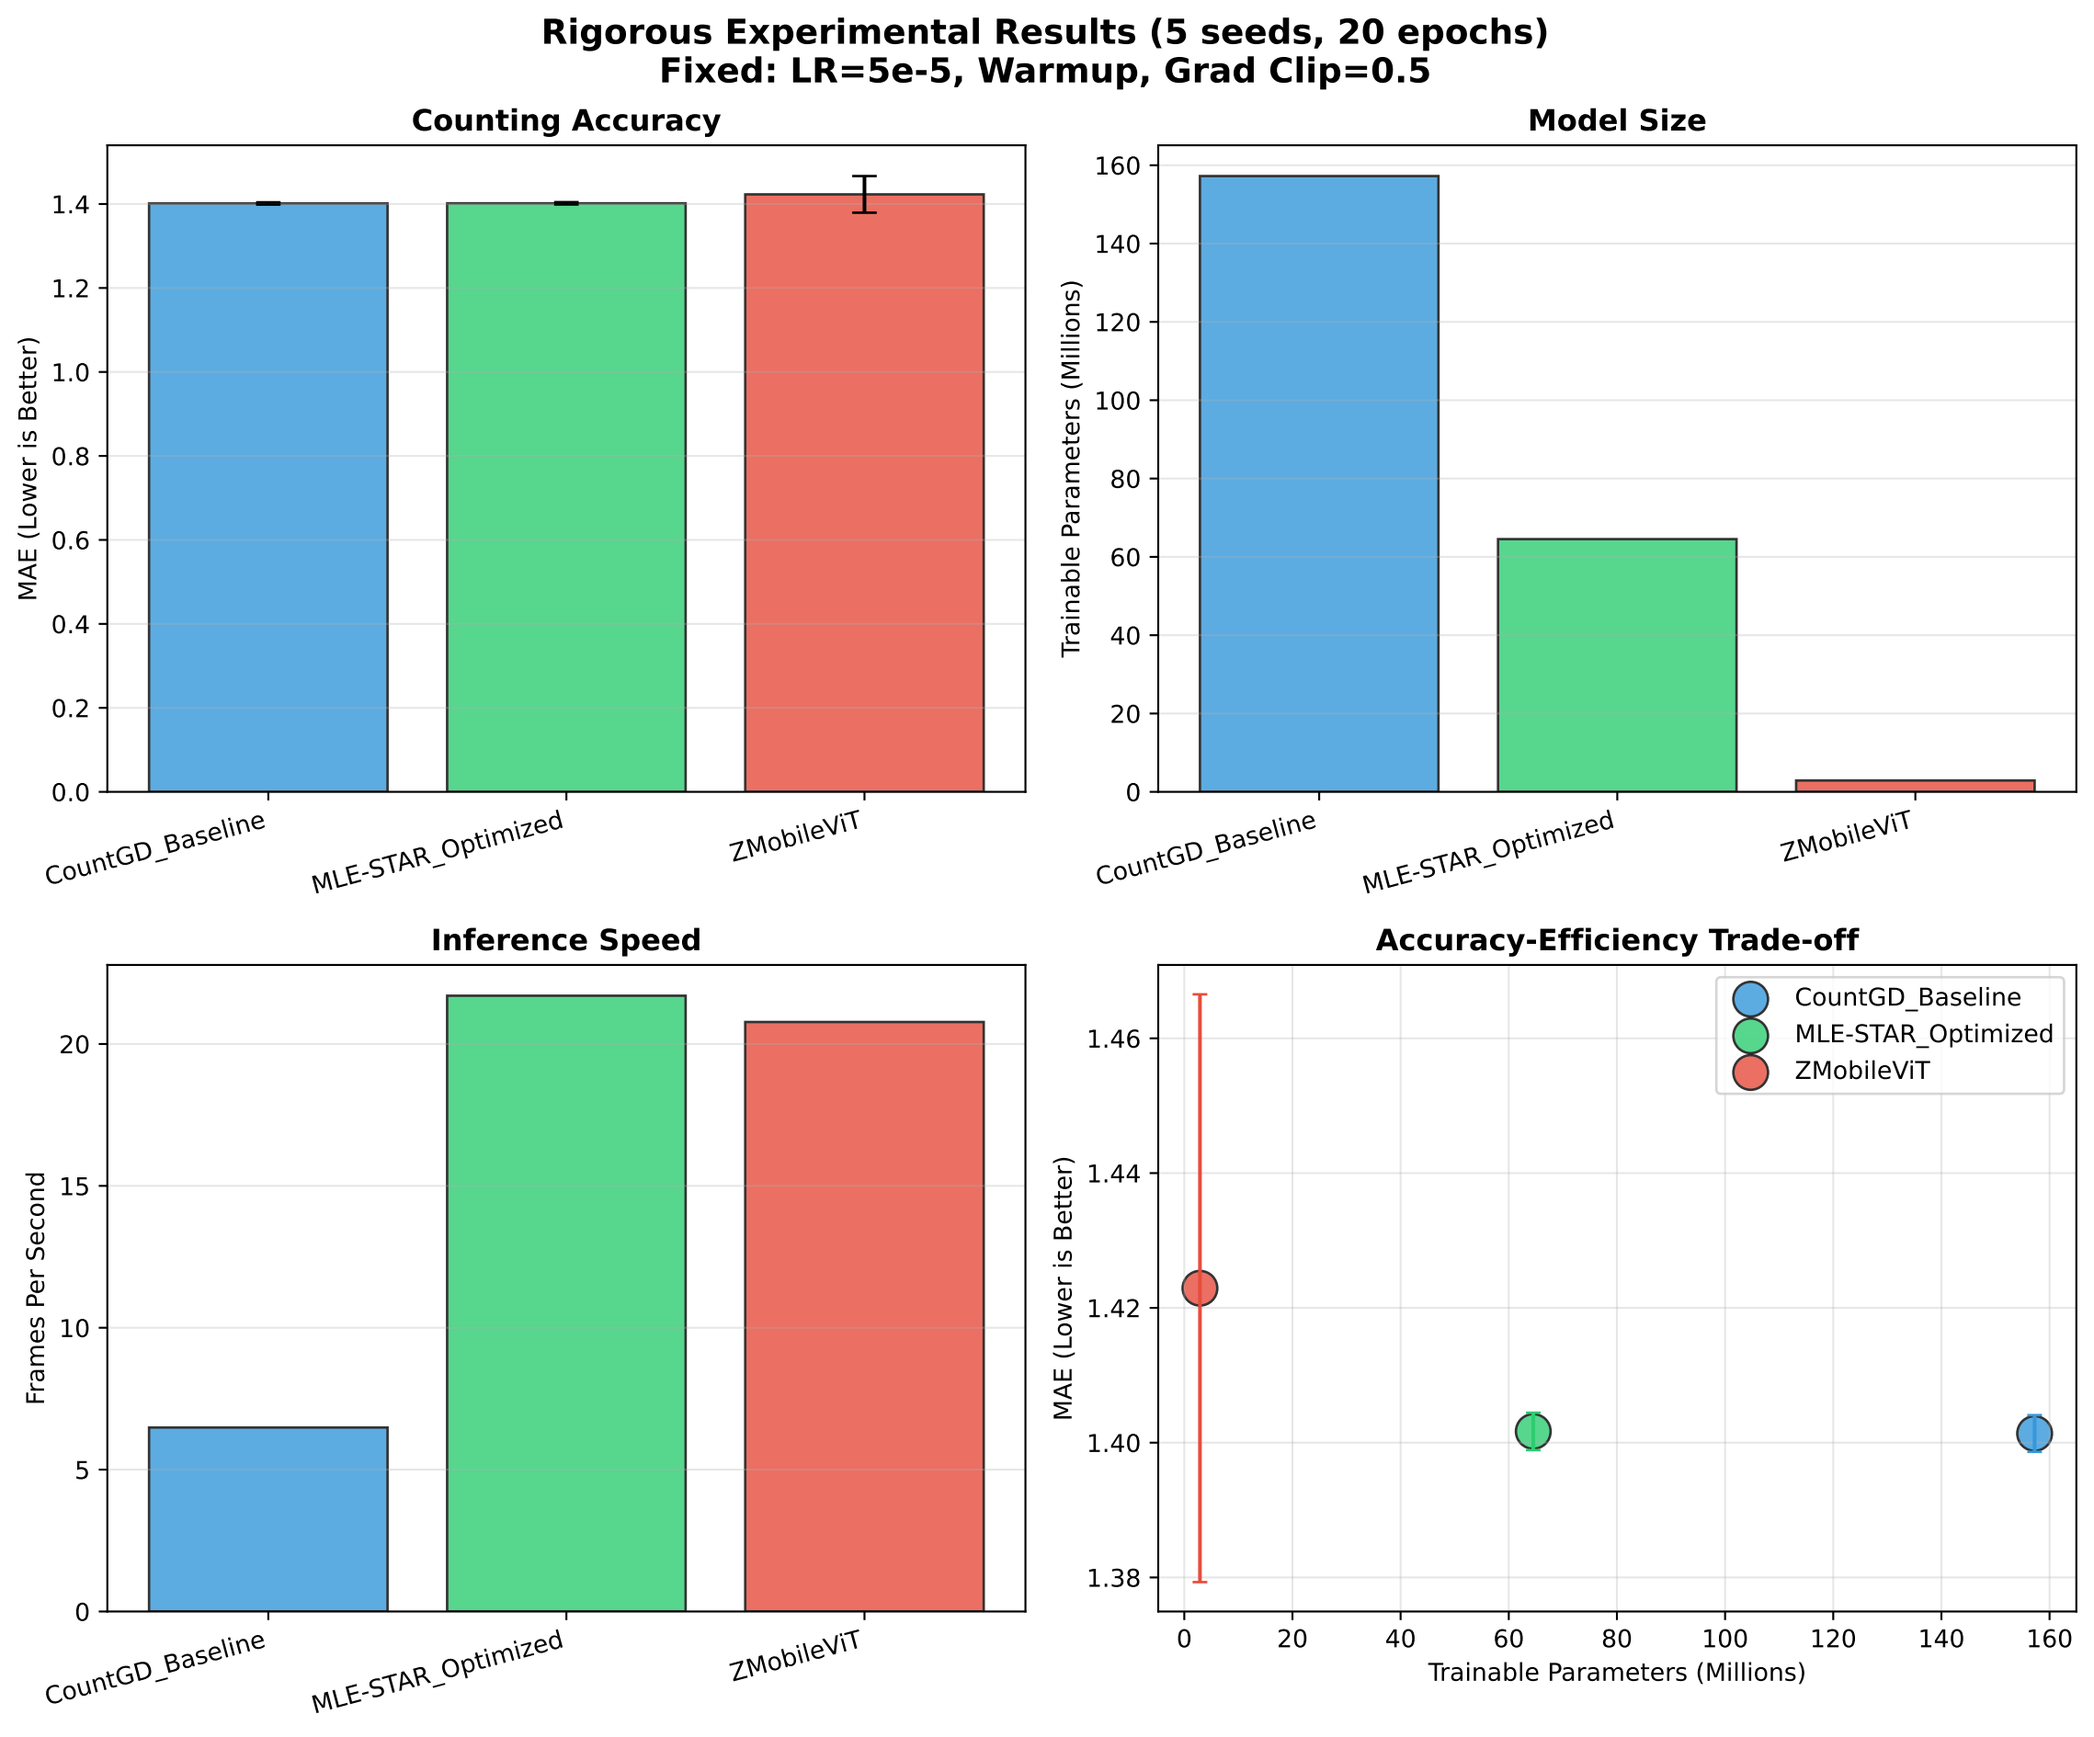
Rigorous Experimental Results (5 seeds, 20 epochs)
Fixed: LR=5e-5, Warmup, Grad Clip=0.5
Counting Accuracy
MAE (Lower is Better)
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
CountGD_Baseline
MLE-STAR_Optimized
ZMobileViT
Model Size
Trainable Parameters (Millions)
0
20
40
60
80
100
120
140
160
CountGD_Baseline
MLE-STAR_Optimized
ZMobileViT
Inference Speed
Frames Per Second
0
5
10
15
20
CountGD_Baseline
MLE-STAR_Optimized
ZMobileViT
Accuracy-Efficiency Trade-off
MAE (Lower is Better)
1.38
1.40
1.42
1.44
1.46
Trainable Parameters (Millions)
0
20
40
60
80
100
120
140
160
CountGD_Baseline
MLE-STAR_Optimized
ZMobileViT